

16720J: Homework 2 - Bag-of-Words for Scene Classification

Name - WENBO ZHAO (WZHAO1)

October 8, 2015

3 Warming up with some theory (10pts)

Question 3.1 (3pts, 2-3 lines)

Given an $N \times N$ image, and an $h \times h$ Gaussian filter, we can convolve the image and the filter in $O(h)$ time by first convolving the image horizontally with an $h \times 1$ filter and then convolving vertically with a $1 \times h$ filter.

Question 3.2 (2pts, 2-3 lines)

No. Because *bag-of-words* takes the features of images (by filtering or feature points extraction) and clusters them into a visual words dictionary. The representation of images is converted to mapping each pixel to the dictionary. After histing the frequency of each visual word in the mapped images, we get totally statistic features, neglecting their spatial distribution information.

Question 3.3 (5 pts, 2-3 lines)

Because we want to use a bank of different filters to extract the local texture information of an image in a sampling-reasonable way, but at the same time depressing the noise as much as possible.

4 Representing the World with Visual Words (40pts)

Question 4.1 (5 pts, 3-4 lines) Theory:

There are 3 groups of filters, 11 in each group. For one group:

The first 3 filters are “*Gaussian*” filters, they smooth the noise in the image.

The second 4 filters are “*LoG*” filters, they smooth the noise and detect the edges in the image.

The third 4 filters are “*Derivative of Gaussians*” filters, they pick up the “*gradient change in a direction*” in the image.

Question 4.2 (15 pts)

Coding question, put your implementation in `baseline/getFilterBankAndDictionary.m`

```
1 function [filterBank,dictionary] = getFilterBankAndDictionary(trainFiles)
2 % This function generates a dictionary given a list of images
3 % - INPUT: * trainFiles: a cell array of strings containing the full path to all images
4 % - OUTPUTS: * filterBank: a cell array of filters from FUNC::createFilterBank
5 %             * dictionary: a visual words dictionary from FUNC::kmeans
6 %
7 % Author: WENBO ZHAO (wzhaol@andrew.cmu.edu)
8 % Date: Oct 3, 2015
9 % Log: (v0.1)-(first draft, written all the functions)-(Oct 3, 2015)
10 %      (v0.2)-(modified: fixed bug: improved: )
11 %
12 %%
13 % debug = 'finished filter response, start k-means'
14 % debug = 'start filter response';
15 %% Set directories
16 % imageDir = '../images/'
17 % switch debug
18 %     case 'start filter response'
19 %% generate filter bank
20 fprintf('Getting filter bank ... \n');
21 filterBank = createFilterBank();
22 fprintf('Done.\n');
23 %% Generate filter responses
24 fprintf('Generating filter responses ... \n');
25 alpha = 100; % [50,150]
26 for i = 1:length(trainFiles)
27     I = imread(trainFiles{i});
28     filterResp = extractFilterResponses(I, filterBank);
29     randPixels = randperm(alpha);
30     filterResp = filterResp(randPixels, :);
31     filterResponses(i,:) = filterResp(:);
32 end
33 filterResponses = reshape(filterResponses, [length(trainFiles)*alpha, 3*size(filterBank,
34 fprintf('saving filter responses ... \n');
35 save('filterResponses', 'filterResponses');
36 fprintf('Done.\n');
37
38 %%     case 'finished filter response, start k-means'
39 % load filterBank.mat
40 % load filterResponses.mat
41 %% Cluster filter response
42 fprintf('Getting dictionary ... \n');
43 K = 200; % [100, 300]
44 [~,dictionary] = kmeans(filterResponses, K, 'EmptyAction', 'drop');
45 fprintf('Saving dictionary ... \n');
46 save('dictionary', 'dictionary');
47 fprintf('Done.\n');
48 % end
49
```

Question 4.3 (5 pts, 3-4 lines) Theory

If the visual words dictionary is too small, it cannot fully represent the filtered responses of images in the training set, thus cannot distinguish the filtered responses well. Contrarily, if the dictionary goes too large, similar responses can be represented by different words, it lacks generalization and is sensitive to noise. It also yields high mapping dimension and exhaustive computation.

Question 4.4 (15 pts)

Coding question, put your implementation in `baseline/getVisualWords.m`

```

1 function [wordMap]=getVisualWords(I,filterBank,dictionary)
2     % This function map each pixel in the image to its closest word in the
3     % dictionary.
4     % - INPUTS: * I: image
5     %               * filterBank: a bank of filters from FUNC::createFilterBank
6     %               * dictionary: a visual words dictionary from FUNC::kmeans
7     % - OUTPUT: * wordMat: an index map of the size(I) that maps each pixel
8     %               response to its closest in the dictionary
9     %
10    % Author: WENBO ZHAO (wzhaol@andrew.cmu.edu)
11    % Date: Oct 3, 2015
12    % Log: (v0.1)-(first draft, written all the functions)-(Oct 3, 2015)
13    %       (v0.2)-(modified: fixed bug: improved: )
14    %
15    fprintf('\nMapping the filtered response of each pixel in image to its closest index in
16
17    % the filtered response for Image
18    % alpha = 100;
19    I_response = extractFilterResponses(I, filterBank); % response of (M pixels * N filters)
20    % compute the distance(ImageResponse, dictionary)
21    wordMap = zeros(size(I_response, 1),1);
22    dist = pdist2(I_response, dictionary);
23    [~, wordMap] = min(dist,[],2);
24    % for i = 1:size(I_response, 1)
25    %     dist = pdist2(I_response(i,:), dictionary);
26    %     [~, wordMap(i)] = min(dist); % find the closest and assign the index
27    % end
28    wordMap = reshape(wordMap, [size(I,1), size(I,2)]);
29
30 end

```

5 Building a Recognition System (95pts)

Question 5.1 (10 pts)

Coding question, put your implementation in `baseline/getImageFeatures.m`

```
1 function [h] = getImageFeatures(wordMap,dictionarySize)
2 % This function returns the histogram of visual words (bag of words) within
3 % the given image.
4 % - INPUTS: * wordMap: an index map of the size(Image) that maps each pixel
5 %             response to its closest in the dictionary
6 %             * dictionarySize: # of visual words in the dictionary
7 % - OUTPUT: * h: a (dictionarySize*1) histogram
8 %
9 % Author: WENBO ZHAO (wzhao1@andrew.cmu.edu)
10 % Date: Oct 4, 2015
11 % Log: (v0.1)-(first draft, written all the functions)-(Oct 4, 2015)
12 %       (v0.2)-(modified: fixed bug: improved: )
13 %
14
15 h = hist(wordMap(:), dictionarySize);
16 h = (h./sum(h)).'; % normalize, sum(h)=1
17
18 end
```

Question 5.2 Extra credit (5 pts 2-3 lines) Theory:

Because by normalizing all histograms by the total number of features in the image, we first make the histograms from different sublayers and subcells comparable, and we also reduce computational cost.

Question 5.3 Extra credit (5 pts 2-3 lines) Theory

Because by weighting different layers inversely proportional to the cell size at the layer, we penalize matches found in larger cells that might induce increasingly dissimilar features.

Question 5.4 (20 pts)

Coding question, put your implementation in `baseline/getImageFeaturesSPM.m`

```
1 function [h] = getImageFeaturesSPM(layerNum, wordMap, dictionarySize)
2 % This function forms a multi-resolution representation of the given image
3 % - INPUTS: * layreNum: # of layers, L+1, [0 1 2 ... L]
4 %             * wordMap: an index map of the size(Image) that maps each pixel
5 %             response to its closest in the dictionary
6 %             * dictionarySize: # of visual words in the dictionary, K
7 %             clusters
8 % - OUTPUT: * h: a ( (K(4^(L+1)-1)/3 * 1) histogram
```

```

9 %
10 % Author: WENBO ZHAO (wzhaol@andrew.cmu.edu)
11 % Date: Oct 4, 2015
12 % Log: (v0.1)-(first draft, written all the functions)-(Oct 4, 2015)
13 %      (v0.2)-(modified: fixed bug: improved: )
14 %
15
16 wordMap = wordMap(:);
17 %% number of cells for each layer
18 numCell = ones(layerNum,1);
19 for i = 1:layerNum
20     numCell(i) = (2^(i-1))^2;
21 end
22 %% index into each cell of the wordMap in each layer
23 index = [];
24 % for i = 1:layerNum
25 %     index(:,i) = [0:length(wordMap)./numCell(i):length(wordMap)]; % e.g.
26 %     % [0 19200 38400 57600 76800] for length(wordMap)=76800 and numCell=4
27 % end
28 index = [0:length(wordMap)./numCell(layerNum):length(wordMap)];
29 %% start from histing the finest layer (L+1)
30 h_2 = [];
31 for j = 1:numCell(layerNum)
32     ind = [floor(index(j)+1) : 1 : floor(index(j+1))]; % index for the jth cell
33 %     h_2 = [h_2; hist(wordMap(ind), dictionarySize).']; % h of size(numCell(layerNum) *
34     h_2 = [h_2; (getImageFeatures(wordMap(ind), dictionarySize)).'];
35 end
36 %% A temporal solution for coarser layers
37 h_2 = reshape(h_2, [numCell(layerNum), dictionarySize]);
38 h_1 = [sum(h_2(1:4,:), 1); sum(h_2(5:8,:), 1); sum(h_2(9:12,:), 1); sum(h_2(13:16,:), 1)];
39 h_1 = bsxfun(@rdivide, h_1, sum(h_1,2)); %
40 h_0 = [sum(h_2(:,,:), 1)];
41 h_0 = h_0./sum(h_0); %
42 %%
43 %% weight
44 weight = 1/2.*ones(layerNum,1);
45 weight(1) = 2^(-(layerNum-1));
46 weight(2) = 2^(-(layerNum-1));
47 for i = 3:layerNum
48     weight(i) = 2^((i-1)-(layerNum-1)-1);
49 end
50 %% sum the weighted hist and normalize
51 h = [weight(1).*h_0; weight(2).*h_1; weight(3).*h_2];
52 h = h(:);
53 h = sqrt(h); %
54 h = h./norm(h,1);
55
56 end

```

Question 5.5 (10 pts)

Coding question, put your implementation in `baseline/distanceToSet.m`

```

1 function [histInter] = distanceToSet(wordHist, histograms)
2 % This function returns the histogram intersection similarity between
3 % wordHist and histograms.
4 % - INPUTS: * wordHist: a (K(4^(L+1)-1)/3 * 1) histogram
5 %           * histogram: a (K(4^(L+1)-1)/3 * T) histogram from T training
6 %           samples
7 % - OUTPUT: histInter: a (1*T) vector of similarity
8 %
9 % Author: WENBO ZHAO (wzhao1@andrew.cmu.edu)
10 % Date: Oct 4, 2015
11 % Log: (v0.1)-(first draft, written all the functions)-(Oct 4, 2015)
12 %      (v0.2)-(modified: fixed bug: improved: )
13 %
14 histInter = sum( bsxfun(@min, wordHist, histograms) ); % check dimensions
15 % must be (K(4^(L+1)-1)/3 * 1) and (K(4^(L+1)-1)/3 * T), histInter is (1*T)
16
17 end

```

Question 5.6 (5 pts)

Coding question, put your implementation in `baseline/createHistograms.m`

```

1 function outputHistograms = createHistograms(dictionarySize, imagePaths, wordMapDir)
2 %code to compute histograms of all images from the visual words
3 %imagePaths: a cell array containing paths of the images
4 %wordMapDir: directory name which contains all the wordmaps
5
6 outputHistograms = []; %create a NumImage x histogram matrix of histograms.
7 %this variable will store all the histograms of training images
8
9 %
10 layerNum = 3;
11 for i = 1:length(imagePaths)
12     fprintf('Outputting histogram %d: %s\n', i, imagePaths{i});
13     wordMap = load(fullfile(wordMapDir, [strrep(imagePaths{i}, '.jpg', '.mat')]));
14     wordMap = wordMap.wordMap;
15     outputHistograms = [outputHistograms, getImageFeaturesSPM(layerNum, wordMap, dictionarySize)];
16 end
17
18 end

```

Question 5.7 (5 pts)

Coding question, provide any helper code you wrote for this section and list the files here. Also make sure your `.mat` is included in your submission.

```

1 % Ishan Misra
2 % CV Fall 2014 - Provided Code

```

```

3 %main script to train your system.
4 %for debugging, comment out function calls that you have run already
5
6 load('traintest.mat');
7 imageDir = '../images'; %where all images are located
8 targetDir = '../wordmap'; %where we will store visual word outputs
9
10
11
12 %%compute filter responses and make dictionary
13 fprintf('Computing dictionary ... \n');
14 computeDictionary(trainImagePaths,imageDir);
15 load('dictionary.mat','dictionary','filterBank');
16 fprintf('done\n');
17
18 %%now compute visual words for each image
19 numCores = 2; %number of parallel jobs to run. Laptops may not support higher numbers
20 fprintf('Computing visual words ... ');
21 batchToVisualWords(trainImagePaths,classnames,filterBank,dictionary,imageDir,targetDir,n
22 fprintf('done\n');
23
24 %%now compute histograms for all training images using visual word files
25 tic
26 trainingHistogramsFile = fullfile(targetDir,'trainingHistograms.mat');
27 dictionarySize = size(dictionary,1);
28 fprintf('Computing histograms ... \n');
29 trainHistograms = createHistograms(dictionarySize,trainImagePaths,targetDir);
30 fprintf('done\n');
31 toc
32 save(trainingHistogramsFile,'trainHistograms');
33 load(trainingHistogramsFile,'trainHistograms');
34
35
36 %%the test code just needs to load trainOutput.mat, so lets store it
37 save('trainOutput.mat','dictionary','filterBank','trainHistograms','trainImageLabels');

```

Question 5.8 (10 pts)

Coding question, put your implementation in `baseline/knnClassify.m`

```

1 function [predictedLabel] = knnClassify(wordHist,trainHistograms,trainingLabels,K, choice
2 %This function
3 % - INPUTS: * wordHist: a (K(4^(L+1)-1)/3 * 1) histogram
4 %           * trainHistograms: a (K(4^(L+1)-1)/3 * T) histogram from T
5 %           training samples
6 %           * trainingLabels: a T*1 labels for training set
7 %           * K: (K) nearest neighbors
8 % - OUTPUT: * predictedLabel: predicted label for new wordHist
9 %
10 % Author: WENBO ZHAO (wzhao1@andrew.cmu.edu)
11 % Date: Oct 5, 2015
12 % Log: (v0.1)-(first draft, written all the functions)-(Oct 5, 2015)

```

```

13 %      (v0.2)-(modified: fixed bug: improved: )
14 %
15 %% find the K nearest neighbors of wordHist in trainHistograms
16 % if ~exist(choiceDistance)
17 %     choiceDistance = 'similarity';
18 % end
19 disp(choiceDistance);
20 % choiceDistance = 'similarity' % might want use PARSE??
21 [~, neighbor] = topKNeighbor(wordHist,trainHistograms, trainingLabels, K, choiceDistance);
22 %% predict the label for wordHist
23 predictedLabel = mode(neighbor);
24 % predictedLabel = predictLabel(neighbor, max(trainingLabels));
25
26 end
27
28 function [distance, neighbor] = topKNeighbor(testSet, trainSet, trainLabel, K, choiceDistance)
29 % This function returns the top K distance of testSet to trainSet and their
30 % labels.
31 %
32 %% choice of distance
33 % default choice
34 % if ~exist(choiceDistance)
35 %     choiceDistance = 'similarity';
36 % end
37
38 switch lower(choiceDistance)
39     case 'similarity'
40         % - similarity
41         dist = distanceToSet(testSet, trainSet); % similarity in this case
42         [dist_sort, pos]= sort(dist, 'descend'); % sort the distance in descending order
43         distance = dist_sort(1:K); % find the top K distances...
44         neighbor = trainLabel(pos(1:K)); % and their corresponding labels in trainSet
45     case 'euclid'
46         % - Euclid distance
47         dist = pdist2(testSet.', trainSet.', 'euclid');
48         [dist_sort, pos]= sort(dist, 'ascend');
49         distance = dist_sort(1:K);
50         neighbor = trainLabel(pos(1:K));
51     case 'cosine'
52         % - Cosine distance
53         dist = pdist2(testSet.', trainSet.', 'cosine');
54         [dist_sort, pos]= sort(dist, 'ascend');
55         distance = dist_sort(1:K);
56         neighbor = trainLabel(pos(1:K));
57     case 'kldivergence'
58         % - K-L divergence
59         KL = @(X, Y) (sum(bsxfun(@times, X, bsxfun(@minus, log2(X), log2(Y)))));
60         dist = pdist2(testSet,trainSet, @(testSet,trainSet) KL(testSet,trainSet));
61         [dist_sort, pos]= sort(dist, 'ascend');
62         distance = dist_sort(1:K);
63         neighbor = trainLabel(pos(1:K));
64     case 'builtinKnn'
65         [idx, dist] = knnsearch(trainSet.',testSet.', 'dist','euclid','k',K);
66         distance = dist;

```



```

67         neighbor = trainLabel(idx);
68     otherwise
69         disp('Undefined distance!\n');
70 end
71 end

```

Question 5.9 (10 pts)

Coding question, put your implementation in `evaluateRecognitionSystem.m`

```

1  %Loading the dictionary, filters and training data
2  numCores=2;
3  imageDir = '../images'; %where all images are located
4  targetDir = '../wordmap'; %where we will store visual word outputs
5  load('traintest.mat');
6  load('trainOutput.mat');
7
8  % Close the pools, if any
9  % try
10 %     fprintf('Closing any pools...\n');
11 %     matlabpool close;
12 % catch ME
13 %     disp(ME.message);
14 % end
15 % fprintf('Will process %d files in parallel to compute visual words ...\n',length(train
16 % fprintf('Starting a pool of workers with %d cores\n', numCores);
17 % matlabpool('local',numCores);
18
19 % predict
20 k = 5; % odd number to avoid tie
21 layerNum = 3;
22 distance = 'similarity'
23 % distance = 'euclid'
24 % distance = 'cosine'
25 % distance = 'kldivergence'
26 % distance = 'builtinKnn'
27 C = zeros(length(unique(trainImageLabels)));
28 tic
29 for i = 1:length(testImagePath)
30     image = imread(fullfile(imageDir, testImagePath{i}));
31     wordMap = getVisualWords(image, filterBank, dictionary);
32     h = getImageFeaturesSPM( layerNum, wordMap, size(dictionary,1));
33     predictedLabel = knnClassify(h, trainHistograms, trainImageLabels, k, distance);
34     fprintf('The %d th image, Label: true: %d, predict: %d \n', i, testImageLabels(i), p
35     C(predictedLabel,testImageLabels(i)) = C(predictedLabel,testImageLabels(i)) + 1;
36 end
37 toc
38 % accuracy
39 fprintf('Confusion matrix:');
40 C
41 accuracy = trace(C)/sum(C(:))

```

Verbatim copy of your confusionMatrix:

K = 1, similarity

24	0	0	0	0	0	1	1	0
0	19	12	2	1	7	1	1	3
0	9	11	1	1	4	0	3	5
1	4	4	28	9	1	7	1	1
2	0	0	5	30	0	5	3	10
1	3	2	2	0	8	0	7	2
0	0	0	2	2	1	19	0	0
0	0	2	0	1	4	0	13	3
1	0	0	0	0	1	1	2	27

K = 5, similarity

25	0	0	2	1	0	0	1	0
0	24	16	2	0	11	1	3	4
0	7	9	1	0	3	0	3	4
1	1	2	26	6	0	10	2	0
2	0	0	6	36	0	3	2	6
0	3	2	0	0	6	0	7	2
1	0	0	3	1	1	19	0	0
0	0	1	0	0	4	0	12	3
0	0	1	0	0	1	1	1	32

K = 11, similarity

27	0	1	0	0	0	0	1	0
0	26	16	1	0	12	1	3	5
0	5	10	1	0	7	0	2	4
0	1	2	29	4	0	12	2	0
2	0	0	7	40	0	4	2	5
0	3	0	0	0	5	0	7	1
0	0	0	2	0	1	16	0	1
0	0	1	0	0	1	0	13	2
0	0	1	0	0	0	1	1	33

K = 15, similarity

27	0	0	0	0	0	0	0	0
0	25	13	1	0	13	2	1	4
0	6	15	0	0	8	0	6	4
0	2	1	29	4	0	10	3	0
2	0	0	8	40	0	6	4	3
0	2	1	0	0	2	0	5	1
0	0	0	2	0	1	16	0	1
0	0	0	0	0	2	0	11	2
0	0	1	0	0	0	0	1	36

```
K = 21, similarity
24    0    0    0    0    0    0    0    0
  0   23   12    2    0   10    2    3    5
  0    5   13    0    1    5    0    4    3
  0    1    4   28    2    0    9    3    0
  5    0    0   10   41    1    5    3    3
  0    6    0    0    0    8    0    5    1
  0    0    0    0    0    1   17    0    1
  0    0    1    0    0    0    0   11    2
  0    0    1    0    0    1    1    2   36
```

Question 5.10 (10 pts)

LAST MINUTE UPDATE: ACCURACY TO 0.6822

Figure 1: LAST MINUTE UPDATE

```
The 317 th image, Label: true: 9, predict: 9

Mapping the filtered response of each pixel in image to its closest index in the visual word dictionary.

similarity
The 318 th image, Label: true: 9, predict: 9

Mapping the filtered response of each pixel in image to its closest index in the visual word dictionary.

similarity
The 319 th image, Label: true: 9, predict: 9

Mapping the filtered response of each pixel in image to its closest index in the visual word dictionary.

similarity
The 320 th image, Label: true: 9, predict: 9

Mapping the filtered response of each pixel in image to its closest index in the visual word dictionary.

similarity
The 321 th image, Label: true: 9, predict: 9
时间已过 307.080723 秒。
Confusion matrix:
C =

    25    0    1    2    1    0    0    2    0
    1   29    3    0    0    9    0    0    2
    0    5   20    0    1    7    1    1    3
    0    0    0   16    1    0    4    0    0
    1    0    0   18   39    0    8    0    1
    0    0    4    0    0    4    0    1    0
    0    0    0    4    0    0   20    0    1
    0    1    2    0    1    5    0   25    3
    2    0    1    0    1    1    1    2   41

accuracy =

0.6822
```

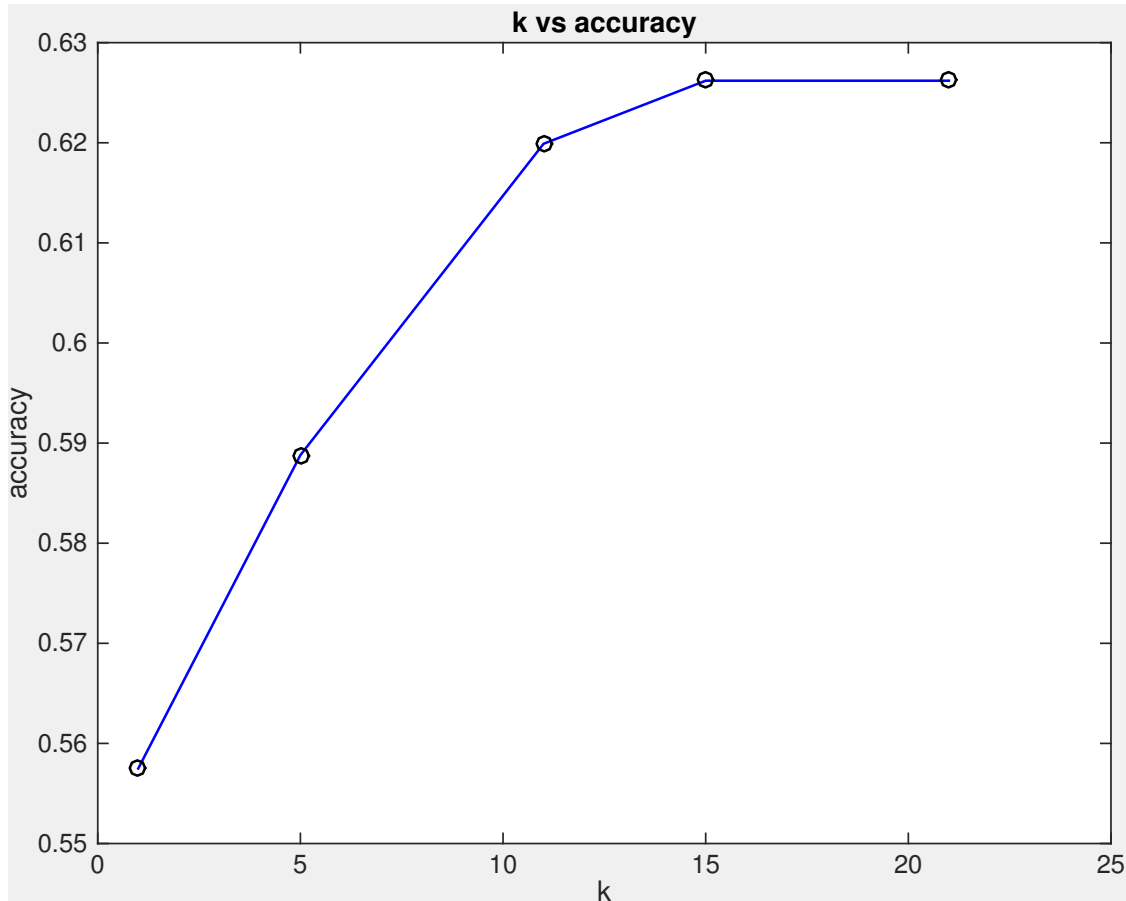
We found when k increases, the accuracy also increases. But when k gets large, the accuracy no longer changes.

A plot of k *v.s.* the classification accuracy (see Fig. 2).

K	1	5	11	15	21	25
Accuracy	0.5575	0.5888	0.6199	0.6262	0.6262	0.6822

Table 1: **Accuracy versus k values.** Add some nice observations here!

Figure 2: k *v.s.* the classification accuracy



Results of classification (see Fig. 3).

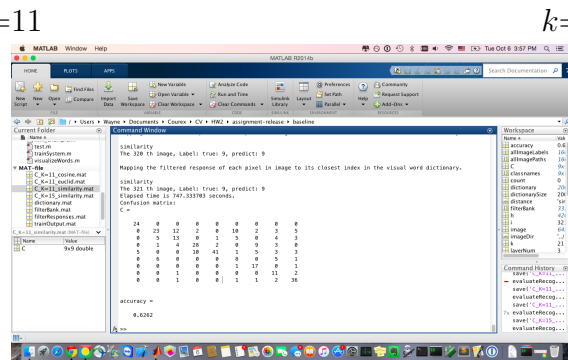
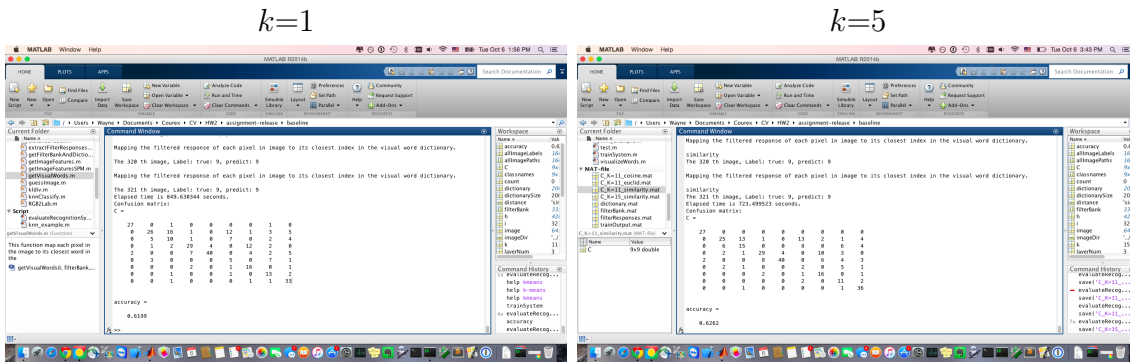
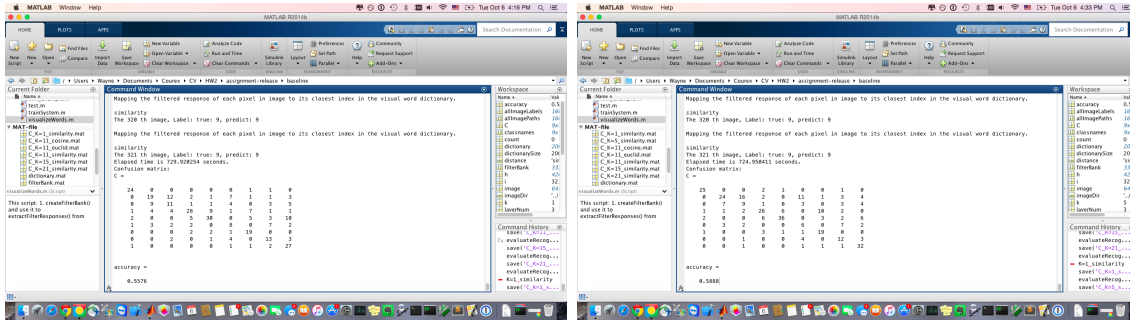
Question 5.11 (5 pts, 2-3 lines): Theory

Nearest neighbor method needs to compute the similarity of data vectors in the dataset. For a large dataset, we need to initialize many centers, and if the vector dimension is high, the computational cost increases. Besides, it does not uniquely weigh different distances.

6 Final thoughts on visual words (40 pts)

Question 6.1 (20 pts)

Coding question, put your implementation in `visualizeWords.m`



$k=21$

Figure 3: Results of classification.

```

1 % This script: finds the 9*9 pixel patch in which the center pixel
2 % corresponds to the index mapping to the visual word in dictionary.
3 %
4 % Author: WENBO ZHAO (wzhao1@andrew.cmu.edu)
5 % Date: Oct 6, 2015
6 % Log: (v0.1)-(first draft, written all the functions)-(Oct 6, 2015)
7 %       (v0.2)-(modified: fixed bug: improved: )
8 %
9 % clear all, close all, clc
10 %
11 % set directories
12 addpath ./export_fig
13 imageDir = '../images';
14 wordMapDir = '../wordmap';
15 % set useful variables
16 load traintest.mat
17 load('dictionary.mat','dictionary');
18 numWords = size(dictionary,1); % number of visual words
19 numPixelInPatch = 9; % numPixelInPatch^2 pixels in a patch
20 % find the pixel patch of visual words in the training image set
21 for i = 1:numWords
22     pixelPatch{i} = zeros(numPixelInPatch, numPixelInPatch, 3);
23 end
24 count = zeros(1,numWords);
25 fprintf('Finding pixel patches for each visual word...\n');
26
27 for i = 1:length(trainImagePaths)
28     image = imread(fullfile(imageDir,trainImagePaths{i}));
29     word = load(fullfile(wordMapDir, [strrep(trainImagePaths{i}, '.jpg', '.mat')]));
30     wordMap = word.wordMap;
31
32     for i=(4+1):(size(wordMap,1)-4)
33         for j=(4+1):(size(wordMap,2)-4)
34             patch = image( (i-4):(i+4), (j-4):(j+4), : ); % read the image patch
35             dictInd = wordMap(i,j); % corresponding visual word index for the current pi
36             pixelPatch{dictInd} = pixelPatch{dictInd} + double(patch);
37             count(dictInd) = count(dictInd) + 1;
38         end
39     end
40
41 end
42
43 % average
44 for i = 1:numWords
45     if pixelPatch{i} > 0
46         avePixelPatch{i} = uint8(pixelPatch{i} ./ count(i));
47     end
48 end
49
50 fprintf('Saving pixel patches...\n');
51 save('pixelPatch', 'pixelPatch');
52 fprintf('Done.\n');
53

```

```
54 % visualize
55 imdisp(avePixelPatch);
```

Add your figure here. See Fig. 4.




Figure 4: **Visualization of words.** The visualization of one word is essentially searching all the images for pixel patches that have their center indexes mapped to this word. The basis idea is that for a pixel in the image that is categorized to a word in the dictionary, its surrounding pixels should at least close to this pixel. So summing up and average all the patches will give us a visual sense of what the word will look like.

Question 6.2 (10 pts)

Add your explanation here with some nice pictures.

Explanation: For the demonstration purpose, I choose three classes from the pool to visualize: bamboo, basilica and railroad. The classification accuracy (as shown in **Question 5.9**) for class *bamboo* and *basilica* is high, while for *railroad* is low, and *railroads* are much wrongly classified as *basilica*. If we want high classification accuracy, we expect the words within class are similar while across classes are distinctive. We can see from Fig. 5, that for *bamboo*, its top 10 words looks similar. But for the latter two their top 10 words looks like they have somewhat unique patterns, especially for *railroad*. Looking in the detail, we found for *railroads* they have more word representations than others. This is due to the image set for *railroad* has many distinctive changes (light, surroundings, etc). These words are of less representation to this class, but might more to classes that share similar words, say, *basilica*. To overcome this problem, we can on the one hand enlarge the dataset or capture new features for this class to extract more generalized patterns, on the other hand reduce the cluster number of this class to avoid sensitiveness to noise.

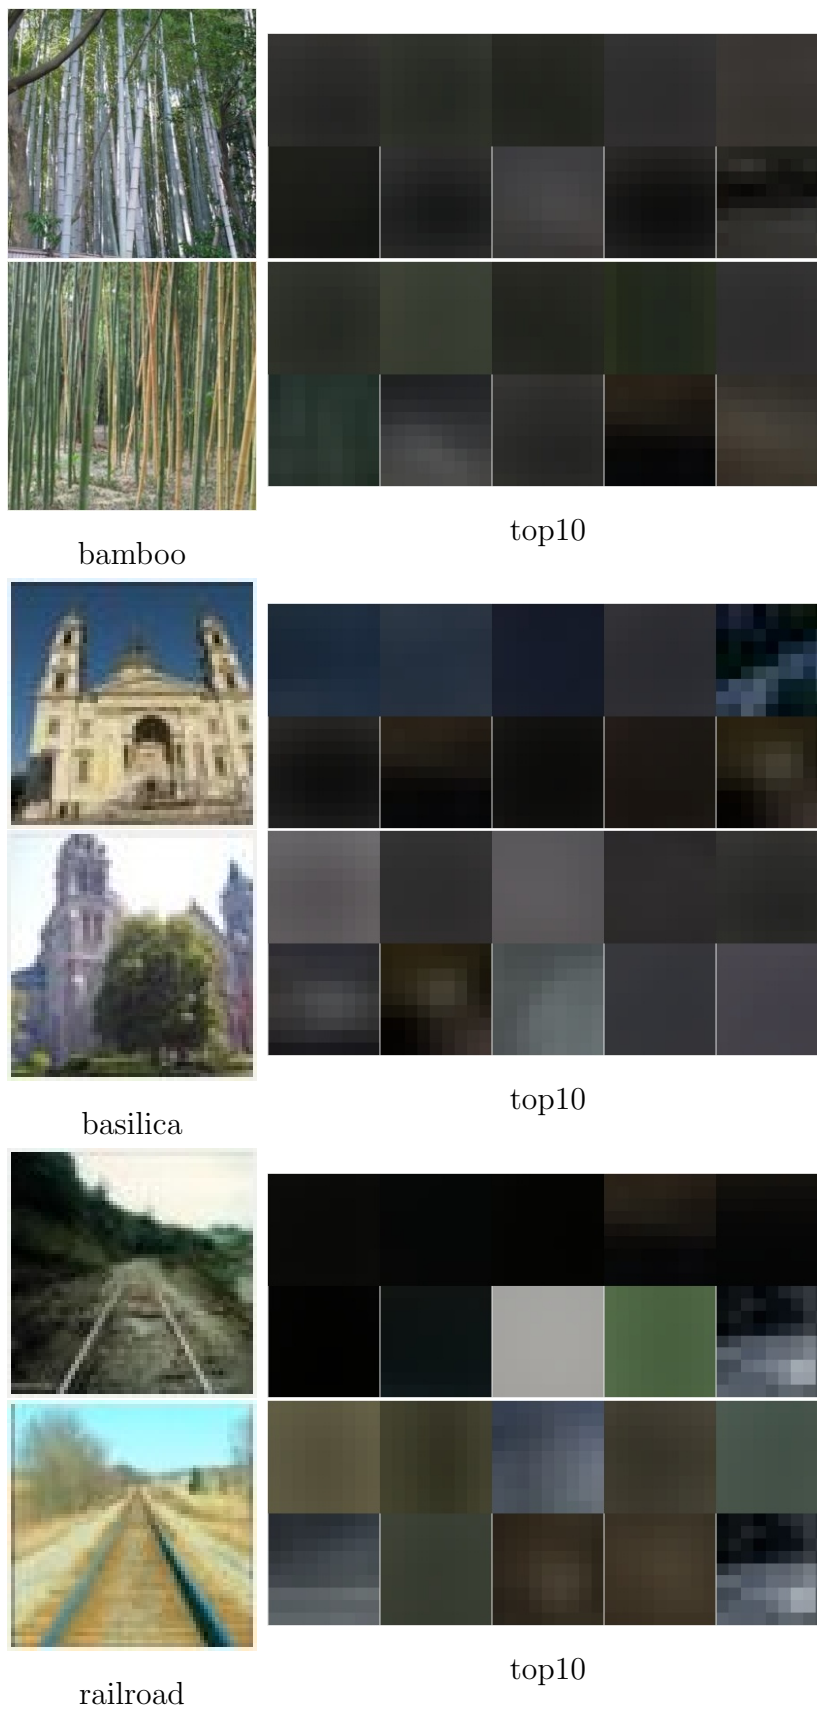


Figure 5: Top 10 words.

Question 6.3 (10 pts, 4-5 lines): Dataset Expansion

1. A way to expand the dataset without using external images is to extract new features from the same dataset. For instance, we can extract interest points features instead of filter-banked features.
2. Flipping will help in bag-of-words with spatial pyramid matching, but not in bag-of-words. Because the former respects spatial information of images, but the latter does not (see **Question 3.2**).

7 Extra Credit: Boosting performance (45 pts)

For the die-hards...

Question 7.1 Better filter bank (10 pts)

The first added is “*Gabor*” filters, they pick up the local spatial and frequency information of images, and sensitive to the edges so that can captures edges of different directions and scales (textures). They are also adaptive to the change of light.

The second added is “*Laplacian*” filters, they pick up the second order derivative of an image isotropically. They are used to detect edges and sharpen the image, but sensitive to noise.

(See Fig. 6.)

Question 7.2 Better image similarity function (5 pts)

I tried *Euclidean* distance and *Cosine* distance for *KNN* classification.


The distance *v.s* accuracy table is shown below in Tab. 2.

Distance	Similarity	Euclidean	Cosine
Accuracy	0.6199	0.4112	0.3988

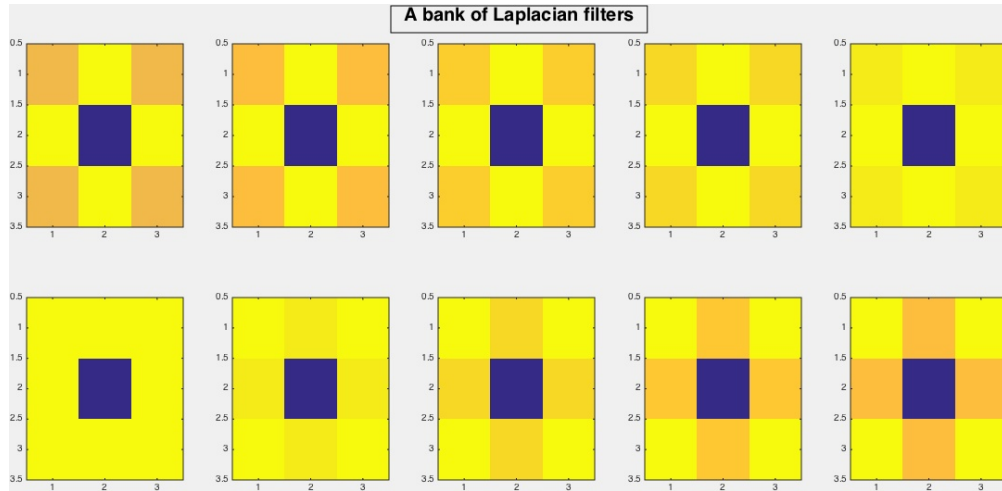
Table 2: **Accuracy versus different distances.** ($k = 11$)

Code for switching from different choice of distances:

```
1 switch lower(choiceDistance)
2     case 'similarity'
3         % - similarity
4         dist = distanceToSet(testSet, trainSet); % similarity in this case
5         [dist_sort, pos]= sort(dist, 'descend'); % sort the distance in descending order
6         distance = dist_sort(1:K); % find the top K distances...
7         neighbor = trainLabel(pos(1:K)); % and their corresponding labels in trainSet
8     case 'euclid'
9         % - Euclid distance
10        dist = pdist2(testSet.', trainSet.', 'euclid');
11        [dist_sort, pos]= sort(dist, 'ascend');
12        distance = dist_sort(1:K);
```



Gabor filter



Laplacian filter

Figure 6: Additional filter banks.

```

13     neighbor = trainLabel(pos(1:K));
14 case 'cosine'
15     % - Cosine distance
16     dist = pdist2(testSet.', trainSet.', 'cosine');
17     [dist_sort, pos]= sort(dist, 'ascend');
18     distance = dist_sort(1:K);
19     neighbor = trainLabel(pos(1:K));
20 case 'kldivergence'
21     % - K-L divergence
22     KL = @(X, Y) (sum(bsxfun(@times, X, bsxfun(@minus, log2(X), log2(Y)))));
23     dist = pdist2(testSet,trainSet, @(testSet,trainSet) KL(testSet,trainSet));
24     [dist_sort, pos]= sort(dist, 'ascend');
25     distance = dist_sort(1:K);
26     neighbor = trainLabel(pos(1:K));
27 case 'builtin_knn'
28     [idx, dist] = knnsearch(trainSet.',testSet.', 'dist','euclid','k',K);
29     distance = dist;
30     neighbor = trainLabel(idx);
31 otherwise
32     disp('Undefined distance!\n');
33 end

```

Question 7.3 Different Features (15 pts)

I tried using SURF descriptors to get the visual word dictionary and mapping images to word maps. The dictionary size is set to 500.

Then the histograms were got as features for classification *WITHOUT* using SPM.

The accuracy is: *0.3396* (see Fig. 7).

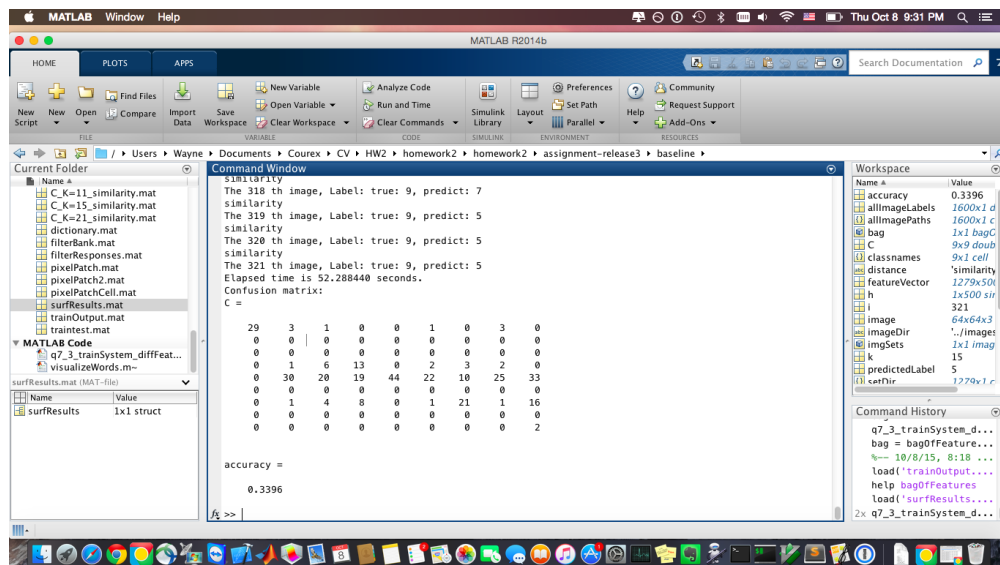


Figure 7: KNN result using SURF-bag-of-words. (K=15)

The code is shown below:

```

1 % This script extracts the visual word dictionary (bag of words) using
2 % feature descriptors, instead of using filter banks.
3 %
4 % Author: WENBO ZHAO (wzhaol@andrew.cmu.edu)
5 % Date: Oct 8, 2015
6 % Log: (v0.1)-(first draft, written all the functions)-(Oct 8, 2015)
7 %       (v0.2)-(modified: fixed bug: improved: )
8 %
9 clear all, close all, clc
10 %% Set directories
11 load('traintest.mat');
12 imageDir = '../images'; %where all images are located
13
14 setDir = fullfile(imageDir,trainImagePaths);
15 imgSets = imageSet(setDir);
16
17 %% Get bag of words using Matlab CV toolbox builtin function
18 % SURF features are used to generate the vocabulary features. Vocabulary is quantized
19 % using K-means algorithm.
20 VocabularySize = 500; % dictionary size
21 PointSelection = 'Grid'; % using SURF feature
22 tic
23 bag = bagOfFeatures(imgSets,'VocabularySize', VocabularySize, 'PointSelection', PointSel
24 toc
25 %% Map image to dictionary
26 % can actually use builtin encode, but...
27 tic;
28 featureVector = encode(bag,imgSet);
29 toc;
30
31 surfResult.bag = bag;
32 surfResult.feats = featureVector;
33 save('surfResult','surfResult');
34 %% knn classification
35 % predict
36 k = 21; % odd number to avoid tie
37 distance = 'similarity'
38 % distance = 'euclid'
39 % distance = 'cosine'
40 % distance = 'kldivergence'
41 % distance = 'builtinKnn'
42 C = zeros(length(unique(trainImageLabels)));
43 tic
44 for i = 1:length(testImagePaths)
45     image = imread(fullfile(imageDir, testImagePaths{i}));
46     h = encode(bag, image);
47     predictedLabel = knnClassify(h.', featureVector.', trainImageLabels, k, distance);
48     fprintf('The %d th image, Label: true: %d, predict: %d \n', i, testImageLabels(i), p
49     C(predictedLabel,testImageLabels(i)) = C(predictedLabel,testImageLabels(i)) + 1;
50 end
51 toc
52 % accuracy
53 fprintf('Confusion matrix:');

```

```

54 C
55 accuracy = trace(C)/sum(C(:))

```

Question 7.4 Encoding (15 pts)

A ‘soft-histogram’ is used to encode the features to visual word dictionary[1], in which:

Instead of assigning one feature to a hard one visual word, a membership between (0,1) is used to measure the belongingness to a certain dictionary.

The membership is computed as:

$$U_k = (1/\|X - d_k\|_2^2) / \sum (1/\|X - d_k\|_2^2)$$

code:

```

1 function [h] = softHist(wordMap, dictionarySize)
2 % This function uses a 'soft-encoding' method to encode the wordmaps to
3 % histograms. Instead of assigning one feature to a hard one visual word,
4 % this function uses a membership between (0,1) to measure the
5 % belongingness to a certain dictionary.
6 % The membership is computed as:
7 %     U_k = (1/||X-d_k||^2_2)/sum(1/||X-d_k||^2_2)
8 % - INPUTS: * wordMap: an index map of the size(Image) that maps each pixel
9 %           response to its closest in the dictionary
10 %          * dictionarySize: K
11 % - OUTPUT: * h: a (dictionarySize*1) histogram
12 %
13 % Author: WENBO ZHAO (wzhaol@andrew.cmu.edu)
14 % Date: Oct 8, 2015
15 % Log: (v0.1)-(first draft, written all the functions)-(Oct 8, 2015)
16 %      (v0.2)-(modified: fixed bug: improved: )
17 %
18 P = pdist2(wordMap(:), [1:1:dictionarySize], 'euclidean');
19 P = 1./(P.^2);
20 P = 1./repmat(sum(P,2), [1, dictionarySize]);
21 h = mean(P,1);
22 end

```

Question 7.5 Getting Fancy! (0 pts)

I can use SVMs to do classification, but since there is 0 pts...

References

- [1] T. Ahonen and M. Pietikäinen, “Soft histograms for local binary patterns,” in *Proceedings of the Finnish signal processing symposium, FINSIG*, vol. 5, p. 1, 2007.